

NaviGator AMS 2018

Daniel Volya, Matthew Griessler, Kevin Allen, Kipling Cohen, Alan Albritton, Nicholas Suhlman, David Zobel, Juan Mejia, Rosemond Fabien, Daniel Olis, Marshall Rawson, Jaxon Brown, Ria Pendon, Marquez Jones, Scarlet Seymour, Dr. Eric Schwartz, Dr. Carl Crane, Dr. Ira Hill, and Shannon Ridgeway

Abstract—NaviGator ASV is a fully autonomous surface vehicle (ASV) built to compete in the Association for Unmanned Vehicle Systems International (AUVSI) Foundation’s 2018 Maritime RobotX Challenge in Oahu, Hawaii. The NaviGator ASV is part of a larger group of collaborative autonomous aerial, surface, and subsurface vehicles known as the NaviGator Autonomous Maritime System (AMS). This paper describes the NaviGator ASV’s structural design, propulsion, power system, electrical design, software infrastructure, and approach to completing the challenges presented in the 2018 Maritime RobotX Challenge.

I. INTRODUCTION

The University of Florida’s (UF) Team NaviGator AMS is a multidisciplinary group composed of undergraduate and graduate students from the departments of Electrical and Computer Engineering and Mechanical Engineering. This project is primarily sponsored by the Machine Intelligence Lab (MIL), which has nearly 20 years of experience in competing in the AUVSI Foundation’s robotics competitions, including numerous championships in the RoboSub and RoboBoat Competitions. Due to the larger scale of the Maritime RobotX Challenge, MIL has partnered with the Center for Intelligent Machines and Robotics (CIMAR), a lab that has competed in three DARPA challenges and has extensive experience with developing highly intelligent large-scale autonomous vehicles. Between MIL’s experience in autonomous maritime systems design and CIMAR’s experience in software architecture design, Team NaviGator AMS puts forth a competitive vehicle for the Maritime RobotX Challenge.

II. VEHICLE DESIGN

This section of the paper will describe the hardware and software that was developed for this competition, as well as the motivations behind these choices. This will include descriptions of early iterations of hardware and software that may have failed, what was learned in that process, and how that knowledge was integrated to improve on the designs.

A. Mechanical Subsystems

The mechanical platform used for the NaviGator ASV is a modified WAM-V research vessel developed by Marine Advanced Research. Several of the mechanical modifications that the team has made will be detailed in this section.

1) *Propulsion*: NaviGator ASV’s propulsion system began as two forward-facing stern thrusters, providing the ASV with a skid-steer configuration. After a short time of testing, it became apparent that adding more thrusters and mounting them at an angle would simplify the vectoring of the thrust to achieve a desired motion, as well as adding the capability of lateral motion. The current configuration features two bow and two stern thrusters oriented at a fixed 45 degrees. This is a thruster configuration that the team used in the 2013 RoboBoat Competition with much success, earning first place. In addition to improved maneuverability, using four thrusters provides redundancy in the system, allowing the ASV to still have maneuverability even if either both bow thrusters or both stern thrusters fail. This feature was invaluable when a motor driver died minutes before a qualification run in the 2013 RoboBoat Competition. With a quick modification to the thruster mapper program, the ASV was able to operate with just three thrusters, saving the run. Moreover, during the 2016 RobotX Maritime Challenge, NaviGator was able to repeatedly and successfully station keep and maneuver the course despite rough currents and winds. The major disadvantage of this configuration is that the fixed angles of the thrusters means that it is not particularly efficient moving in any direction. However, as was demonstrated in 2016, for the tasks that the Navigator ASV is designed to perform, maneuverability provides a significant advantage in maintaining stability.

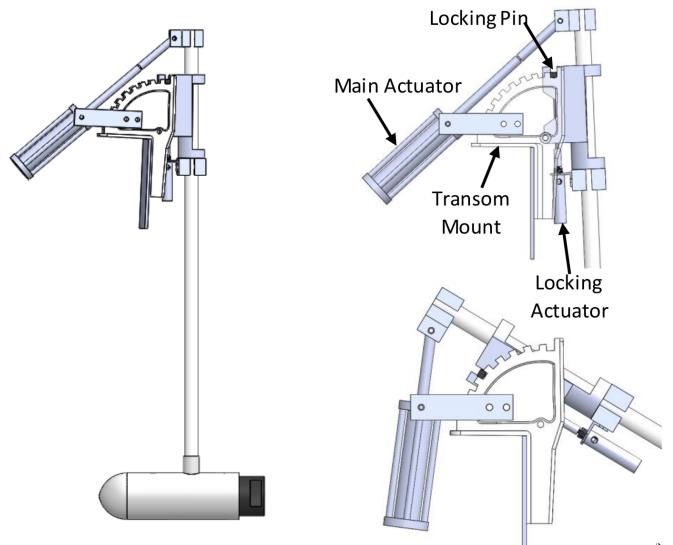


Figure 1. Auto-deploy system for thrusters

Mounting the thrusters posed many challenges and required several design iterations, especially for the bow thrusters. For the ASV to be deployed from a trailer, the bow thrusters had to be either removed or raised during deployment so they would not collide with the trailer structure. Thus, we implemented a student designed auto-deploy system for the thrusters (Fig. 1), which also aided in the efficiency of deploying NaviGator. The system uses two pneumatic linear actuators to rotate the thruster about the transom mount and to lock the thruster in place. The main actuator is mounted in a tandem style to allow it to pivot as the system rotates. To retract the thruster, the locking actuator extends and moves the spring-loaded locking pin. The main actuator then retracts pulling the thruster to an upward position. The air pressure to the locking actuator is then released causing the locking pin to fall into a lock position. The air pressure is then released from the main actuator as it no longer needed to hold the thruster in place. To deploy the thruster to do the same process but in reverse. This system not only saves time but also makes it so that only one person has to get in the water to pull the boat onto or off of the shore.

2) Sensor Mast: The need for a stable sensor platform is paramount in machine vision applications. The preliminary design utilized an 80/20 aluminum rail truss, which did not provide the required stiffness and resulted in smearing of the vessel's detection data. The initial sensor platform also did not raise the LIDAR system high enough to permit detection of obstacles in immediate proximity to the pontoons, a problem rectified in the final design.

As previously mentioned, the cameras, LIDAR, and GPS antenna require a rigid support. The need for an unobstructed GPS antenna guided the design towards a mast structure. For transport to the competition site, the assembly had to fit within the prescribed envelope of a Pelican Products transport case, requiring a modular assembly process. These target specifications led to a base-and-tree assembly, where the mast is simply welded to a plate that then fastens to the payload tray via a superstructure. For corrosion resistance and manufacturability, 6063 aluminum was chosen. To simplify the assembly process, fastener types were standardized. The mast is centered laterally on the ASV, which helps create a well-defined coordinate system that permits simpler software transformations. The sensor mast can be seen in Fig 2.

3) Electronics Box: NaviGator ASV's electronics are housed in a Thule Sidekick cargo box. The team originally considered commercial waterproof boxes, but began looking for other options due to their high costs. One student suggested the idea of using a cargo box after being inspired by family road trips they had taken when they were younger. While traditionally used to mount on the top of cars to provide additional storage, the cargo box was an ideal electronics enclosure due to its watertight integrity, aerodynamic form factor, low cost, and a side-opening mechanism that makes it very easy to access all of the

electronic components. The box's watertight integrity prevented the team from using air circulation for cooling. Instead, a combination of techniques are used to cool the box. First, an adhesive reflective covering was applied to the lid of the box to reflect heat generated by solar radiation. Second, the box has an active water cooling system that is used to remove the heat generated from the electronic components inside the box. Fiberglass inserts were used to mount all of the components inside of the box. These inserts add rigidity to the relatively flimsy box and make it easy to add or remove components from the box. The components that need to be frequently removed, e.g., the hard drives, are attached to the fiberglass with Velcro. The rest of the components are attached with traditional fasteners.

4) Racquetball Launcher: A system for delivering the racquetballs into the target for the Detect and Deliver task was developed by breaking the challenge into subtasks that were solved independently. The two main subtasks that were considered were moving the balls into the target and feeding the balls to the mover. Several ideas for moving the balls were considered, ranging from a catapult to a robotic arm that would drop the balls into the target. Prototypes of several designs were built and tested. One design featured two counter-rotating wheels attached to the trolling motors that were once part of PropaGator 1, the team's submission to the 2013 RoboBoat Competition. The trolling motors were originally used as part of an early prototype, but since they were already waterproof, were effective at launching the balls consistently, and were readily available, the trolling motors were incorporated into the final design. The ball launching mechanism was designed so that any type of ball feeder could be integrated into it. This allowed the team to test multiple types of ball feeder mechanisms, including a carousel and a linear actuator. A prototype of the linear actuator racquetball launcher can be seen in Fig 3. After the carousel mechanism was found to be prone to jamming, the linear actuator design was selected. Additional design criteria that were considered were the ease of loading balls and how quickly all four balls could be launched. Ball loading was addressed by designing a spring-loaded 3D printed ball magazine that is easily detachable. To achieve rapid-fire, the team developed a closed-bolt system. After a ball is loaded into the chamber, it is withheld at the minimum distance from the wheels to reduce the amount of time it takes to fire. In order to prevent premature firing, a retention lip is used.

B. Electrical System

The design goal for the 2016 NaviGator platform was robustness and simplicity. Having achieved those goals for the 2018 competition the design goal was to improve upon the 2016 system.

1) Power Distribution: The 2018 NaviGator retains the dual battery power supply and the power merge board from

the 2016 system. NaviGator remains MIL's biggest project in terms of power required. In 2016 the team designed a power system using two Torqueedo Power 26-104 batteries. The 2016 NaviGator did not provide any feedback on how much current was being drawn from each battery at a given point in time and the voltage level sensing of each battery could be improved. In 2018 NaviGator improved upon the 2016 power system to provide moment to moment, accurate, current and voltage sensing. After the current and voltage sense device there are two power paths that extend from each battery. The high power path connects each battery to two thrusters through fuses. The low power path connects each battery to NaviGator's sense and compute devices through the power merge board.

2) Power merge board: The power merge board is a student-designed printed circuit board assembly (PCBA). It uses Texas Instruments LM5050-2 ideal diode controllers to balance and parallel the two batteries into a single 24V rail to power NaviGator's sense and compute devices. The main benefits derived from the power merge board are two-fold. The everyday benefit is that NaviGator's batteries can be switched out without the computer and networking equipment turning off. The fault tolerant benefit is that if a battery fails, NaviGator retains control and can be commanded. One of MIL's strengths is that parts of MIL's vehicles are designed to be translatable to other vehicles in the lab. This is the third vehicle that has used this power merge board design. The PCBA was designed for PropaGator 1 and then used on PropaGator 2. Both PropaGators have competed in the RoboBoat competition.

3) Passive sonar: The ability to track a point source of sound in the water is encapsulated into the passive sonar pressure vessel. It contains a passive sonar amplification and filtering board (Fig. 2), necessary power regulation, and USB communication. An Analog Digital 4-channel Data Acquisition ADC (ADAR7251) is used to simultaneously sample, amplify, convert, and filter the four incoming signals. The board was designed by Sylphase – a startup founded and run by a former MIL student – and is capable of tracking multiple acoustic sources simultaneously, provided they are at different frequencies.

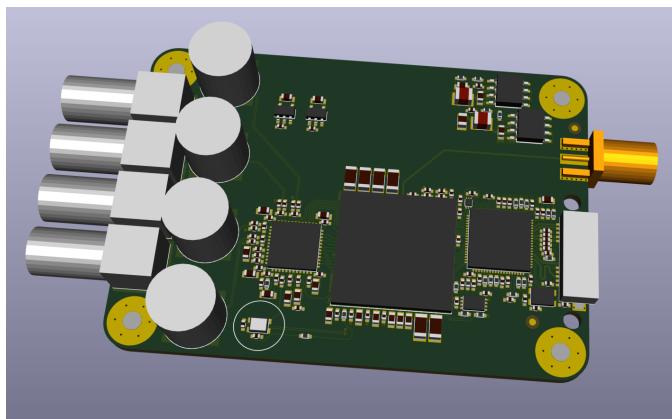


Figure 2. Passive Sonar PCB

4) Kill system: In accordance with the RobotX Kill Switch Specifications, the NaviGator ASV disconnects power to its' thrusters through an emergency kill system capable of operating independently of the motherboard. Power supplied into the motor controllers first passes through four parallel F7 Series power relays, which are controlled by the Emergency Kill Board. This Kill Board receives power from a 22.2V LiPo battery, independent of the vehicle's main batteries. A microcontroller monitors the status of the four E-Stop buttons on the WAM-V's four arms, and deactivates the power relays when these are pressed. Additionally, the Kill Board communicates with the motherboard through a USB connection, constantly relaying the kill switch status and receiving an ongoing "heartbeat" message. Should the motherboard stop sending this heartbeat (indicating software failure) for longer than 5 seconds, the Kill Board will cut power to the thrusters. The Kill Board is also equipped with a Linx NT Series RF transceiver, constantly communicating with another transceiver on the Emergency Controller on a 903.37 MHz carrier frequency. As this transceiver independent of the Wi-Fi connection through the Ubiquiti antenna, the Emergency Controller can kill the power to the thrusters even when the vehicle loses connection to shore. Further, the controller can be used set NaviGator to "emergency control" mode, allowing the user to pilot the vehicle if recovery via shore controls is not available. To cut power, the kill board opens the contacts on the four relays connected to the power on the four motor controllers, cutting power to NaviGator's actuation systems while the computer remains active. The kill board is also used to control NaviGator's indicator lights.

5) Siren System: The siren board is a student-designed PCBA that communicates with the computer over the on board, low speed, CAN network. It controls the siren that wards off curious watercraft during testing of NaviGator.

6) Current and Voltage Sensing Board: The current and voltage sense board is a student-designed PCBA that senses the voltage level and current being drawn by each battery moment to moment. The sensed data is sent to the computer over the on board, low speed, CAN network. There is one PCBA installed in line with each battery. By recording the current drawn from each battery and voltage level over time the health of the batteries can be ascertained both for long term battery health and short term determination of when the batteries need to be switched out and charged.

C. Software System

1) Object Detection and Classification: The lowest level perception service available on the NaviGator ASV is the Occupancy Grid Server. Occupancy grids are a two dimensional grid-like representation of the environment generated by the sensor suite available on the ASV. The generated map contains both the occupied and unoccupied regions in the environment. This information is provided

to the server via any range-detecting sensor onboard. On the ASV, the primary range-detecting sensor is a Velodyne VLP16 LIDAR. A LIDAR uses lasers to provide relatively dense range information of the environment. This information is then segmented by regions containing dense clusters of relatively close points. These bounding regions are treated as obstacles, and are placed in the occupancy grid. This information is then provided to higher level services such as the motion planner and Classification Server. In the Classification Server, the points generated by the LIDAR are clustered into regions on the occupancy grid where it decides which of these distinct regions are objects. The ASV then looks at the bounding box of this object and classifies the object based on the dimensions of its bounding box. The software detects if the object has a prominent plane. If it does, then this information is attached to the object. These objects are then accessible to other programs through the use of a list of detected objects.

2) Motion Planning: For a safer and more flexible planner, the team sought out an algorithm that can handle strict, well-defined constraints. The rapidly-exploring random tree (RRT) algorithm is highly efficient for this scenario [1]. The algorithm starts with a seed node at the ASV's initial state. It then randomly samples a state in the region of navigational interest. A nearness function is applied to every node currently in the tree, and then that node is extended or steered towards the random state following a policy function. The endpoint of that extension is added as a new node to the tree only if it is allowable, and the algorithm repeats. If an extension, or any intermediate state leading up to it, is not allowable, that iteration is simply abandoned. Once a node reaches the goal region, the tree is efficiently climbed from the goal back to the seed, and is classified as one solution to the planning problem. The best of the found solutions is defined as the one that takes the least amount of time. The goal region is likely to be reached because one can bias tree-growth towards it by shaping the probability density function from which random states are sampled.

After selecting the RRT algorithm for safety and flexibility, the final step was to integrate the algorithm with a real-time system. One of the biggest difficulties in doing this was dealing with a highly nonstatic environment. Obstacles spontaneously appear when they get in range of the perception system. This means that a valid path can suddenly become invalid with only seconds to spare. To make efficient use of time, the planner should always be planning the next move so that the RRT has more time to get a better solution. To handle this, the planner had to be made asynchronously interruptible, and a lot of planreevaluation and crisis-aversion logic had to be built in to elegantly deal with spontaneously appearing and/or moving obstacles that cross the ASV's current path. The ASV's real-time ROS-integrated RRT algorithm being run for an arbitrarily drawn, complicated occupancy grid can be seen in Fig 4. Tree nodes can be seen in blue. The ASV

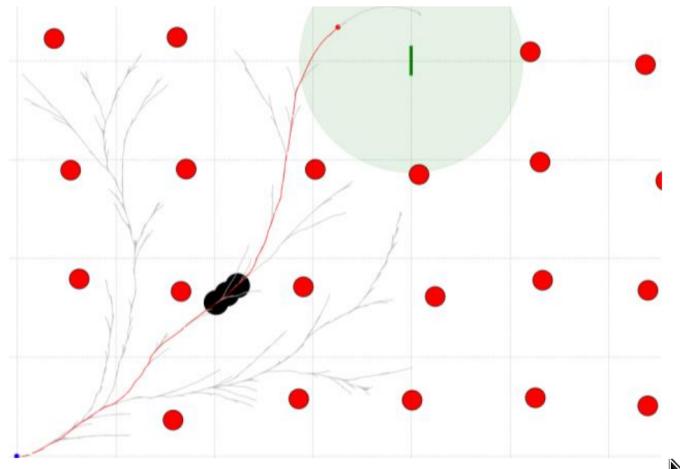


Figure 3. An example of the NaviGator ASV's RRT planning towards a goal region

was only given one second to plan its first move. It used its time during the first move to plan its second move, shown in red. While the paths generated using this method are safe and useful for solving the problem of navigation in the competition, the team is actively working on improved heuristics for smoothing out the paths.



Figure 4. The NaviGator ASV's real-time ROS-integrated RRT algorithm being run on an arbitrarily drawn, complicated occupancy grid

3) Motion Control: Since the RRT motion planner uses a model of the ASV, in principle it would be possible to employ a model-predictive control architecture in which the ASV rapidly re-plans from its current state to steer it back onto the desired path. However, due to the randomness inherent to the RRT itself, such a method did not work well in practice. Thus, the team opted to make use of the sequence of states generated by the motion planner rather than the inputs to define the reference a feedback controller tracks. First, a simple manually-tuned full-state feedback PD controller was used. Tracking along straight paths was nearly perfect with this alone, providing a positional steady-state error of less than 0.25 meters. However, along curves, a larger positional steady-state error of a few meters would always emerge depending

on the curvature. Even the introduction of a standard integral term did not fix this problem. The team figured that this was because an integral of the world-frame error alone would only be able to compensate for disturbances that are constant in the world-frame. Simulation revealed that the sources of the curved motion disturbances were centripetal-Coriolis effects and heading-dependent drag forces. A more intelligent integrator would be necessary to compensate for these state-dependent disturbances. Most marine and aerial systems accomplish this by using a model-reference adaptive control (MRAC) architecture. A block diagram of the MRAC controller used for the ASV is shown in Fig 5. In this diagram, y_{ref} is the current state in the sequence generated by the motion planner, u is the control effort choice, and y is the actual state. MRAC works very well on the ASV, bringing steady-state error to negligible amounts in all cases without introducing oscillations. Additionally, it does not wind-up as much as an ordinary integrator when unexpected disturbances are applied, such as humans pushing the ASV, since it is trying to adapt specifically to drag and inertial effects instead of constant external forces. Finally, with the controller outputting desired wrenches, the last operation needed is to map that wrench to a thrust command for each thruster. A surface vehicle would only need three thrusters to be holonomic, but with four, the ASV is more fault tolerant. This redundancy in the mapping can be solved as a regularized least-squares problem by evaluating a pseudoinverse [2].

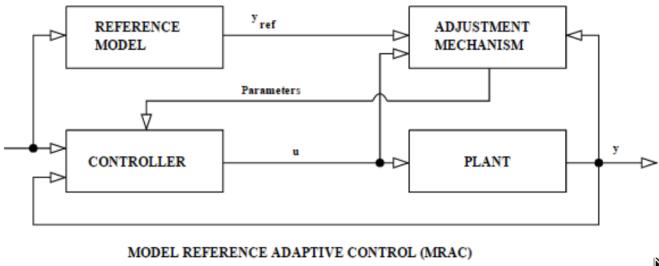


Figure 5. Block diagram of the MRAC controller used on the NaviGator ASV [3]

4) Navigation and Odometry: The NaviGator ASV uses a student-developed Sylphase global positioning system (GPS) and inertial navigation system (INS) that is in the process of being commercialized by Forrest Voight, a UF student and member of Team NaviGator AMS. It primarily consists of a circuit board with a Spartan-6 field programmable gate array (FPGA), radio frequency (RF) frontend, inertial measurement unit (IMU), magnetometer, and a barometer. The FPGA performs the correlation operations that enable tracking of GPS satellites. All the sensor measurements and correlations are passed to a computer via USB, into a pipeline of software modules that track and decode the signals from the GPS satellites and then fuse measurements using an extended Kalman filter into an estimate of the ASV's pose in both absolute world and relative odometry coordinate frames. Last, the

resulting odometry is transformed so that it describes the ASV's coordinate frame and it is then passed to ROS. By using the sensors to aid the GPS solution and taking advantage of GPS carrier phase measurements, extremely precise relative odometry is possible, with noise on the order of centimeters over periods of seconds to minutes. This is the result of years of work, during which several iterations of the hardware were produced. The initial version of the hardware was a Beaglebone cape, but quickly moved to the USB/FPGA approach for ease of development and reduced CPU load. The current revision of the hardware is shown in Fig 9.

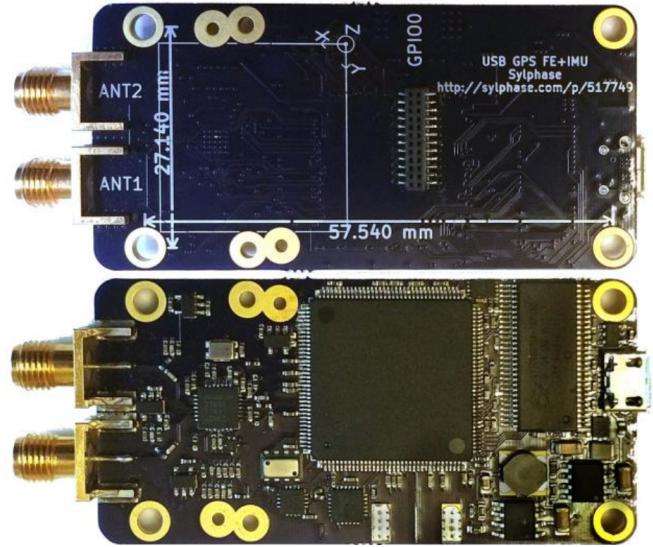


Figure 6. Current hardware revision of the Sylphase, a student-designed GPS/INS

5) State Machine: The state machine that is used in solving the challenges uses a directed acyclic graph (DAG) to decide which missions to complete at which time. Each mission is first defined by three key attributes: the other missions that it depends on, the objects that it depends on, and whether or not the mission should be re-executed. For example, the Scan the Code challenge does not depend on any other challenges, it depends on the Scan the Code object being recognized after it is executed, it should not be re-executed after it is completed. The state machine is constantly listening for new objects to be found. Once one is found, it goes to the parent missions in the DAG and evaluates if they are ready to be completed. If one of these missions is ready, it is executed. Once it is complete, the DAG is reevaluated for more missions to be complete. This continues until all missions are complete.

III. DESIGN STRATEGY

One of the most difficult tasks in developing autonomous vehicles is the detection and recognition of objects, which is then passed down to higher level decision planning. In the past, our team would develop custom and traditional computer vision algorithms for object recognition, however

this requires hours of development time, and often results in complex solutions with mediocre reliability. As such, we have integrated deep neural networks as an initial guess framework, and developed a pipeline to quickly train and test the network. This has resulted in a faster debugging process, and created a central framework for which many design decisions evolved from.

A. Deep Neural Networks

Machine learning has become one of the integral components for perception solutions on all our projects. The ability of machine learning to quickly give us a targeted region of interest without having to craft a traditional computer vision solution has drastically cut down on development time for our systems. Due to our increasing usage of neural networks, efforts were made to develop a fully featured development pipeline for the purpose of training and deploying neural networks for computer vision or perception related tasks. To accomplish this, we employed the Tensorflow Object Detection library which was custom compiled to work with CUDA 10 and the new NVIDIA tensorcore architecture. We used the Labelbox labelling tool to handle all of the manual data processing.

1) Data Handling: One of the well known drawbacks of deep neural networks is the tremendous data requirements for achieving any semblance of accuracy in object detection. To combat this, we utilized a tactic called transfer learning. Transfer learning is the process of taking a network that was trained on a separate dataset (for our purposes this was most commonly the COCO dataset) and retraining the final layers of the network on our own datasets. This takes advantage of the fact that the majority of the neural network is taken up by general shape and color differentiation. The majority of the data requirement is due to these early and middle network layers being trained to differentiate shapes. Once the shapes are learned, only a small amount of data is required for learning the finer details. It is only in the final layers of a neural network that the finer details of an object are discerned and analyzed. Thus we targeted these layers for retraining and held the other layers to be constant. This reduced the number of training images required for each network down from potentially tens of thousands of images to a couple hundred.

With this in mind, during our weekly testing days, we recorded footage of design objects with the mounted cameras on NaviGator from many different angles and in a myriad of lighting and weather conditions. The ROS bags containing this camera footage was processed and segmented out, so anyone who wishes can access and download the images generated from them. This is publicly available for any team to use if they find it useful, as is the code for the pipeline. Note that at this stage the data is not labelled however. Deep neural networks require that we have ground truth labels in order for the network to actually learn anything. We employed Labelbox for this

purpose, as it allowed for collaborative labelling, so multiple members of the team could process the same dataset. This sped up the process considerably.

Once the data is labelled it must be downloaded and processed into a format that Tensorflow recognizes and can use. Additionally due to flaws in the Labelbox software, some of the labels could extend outside of the image bounds or be the size of a single pixel. These ‘broken’ labels could seriously hinder the network’s ability to train off the dataset or cause the training process to crash entirely. Thus arose the MLL Machine Learning Pipeline.

2) Training Networks: The pipeline was developed using python scripts, docker containers, and a couple bash scripts. The central premise of the pipeline is to download the images directly from Labelbox using the JSON file that can be exported from Labelbox. The images and labels are downloaded as png and xml files respectively. These files are then separated into a 60 40 split for training and testing data respectively. Once broken up, we generate two separate CSV files that arrange the labels into the format required by Tensorflow. At this stage we also perform the error checking on the bounds. We check to ensure that the labels are larger than a given area we define and that the labels do not exceed the bounds of the image itself. Spelling check and label validity are also checked against the specific project needs at this stage. For example, if we are training a network to analyze totems and buoys, we will toss out labels and images only containing the docks. Once the data is processed, it is compiled by Tensorflow into TFRecords. These records combine both the images and the labels into a binary file that can be handled more easily by Tensorflow itself. This is what is loaded into the actual training script. The files we generate previously are then cleaned up automatically if desired. If not the generated files are left and can be used to validate the integrity of the dataset through a separate script which visualizes the labels using OpenCV. At this stage the user will select a pretrained model that fits their needs. We found the COCO dataset to be adequate for our purposes and downloaded the architecture that had a good balance of accuracy and speed, as we require real time object identification. There are plenty of options available at the Tensorflow Model Zoo, part of their Object Detection repository on github. After selecting a model, you can need to make some edits to its configuration file so that the model knows where to load its training and testing data from. There are a variety of other options that can be useful to change and tweak to give better accuracy but this is dependent on the model choice. The docker image now comes into play. Launching the pipeline script we created spins up a docker image that contains all the necessary software prerequisites for Tensorflow and compiles it from source. This avoids requiring the user of our pipeline to download the repository locally and install the myriad of dependencies required therein. With this done one can easily train any network from the model zoo and with any dataset they desire. The output will be a fully trained network with a frozen

inference graph that can be used on any system running a compatible version of Tensorflow.

3) Perception Application: Now that we have a trained network, we employ a manager that spins up processes. Each process is devoted to a specific challenge, but the code within is in essence the same. The process will load the network associated with the challenge and begin processing images fed to it from our cameras. It will then publish a bounding box corresponding to its observations. We can set thresholds for confidence levels and size of bounding boxes to further refine our results from this stage and apply more traditional computer vision techniques specific to each challenge, but the bulk of the work is now complete. Machine learning has greatly enhanced and streamlined our solutions to computer vision challenges. All work we have done is available for others to use and modify, and we encourage other teams to explore our process to develop a more robust pipeline and networks to solve the computer vision challenges.



Figure 7. Demonstration of classification using a deep neural network

B. Identify Symbol and Dock Challenge

One particular example that demonstrates the trade-off between reliability and development time is in the Identify Symbol and Dock Challenge. This mission begins by selecting the dock object from the object server, which is easily identified by being the largest connected object on the course. The object server gives the mission a rotated box enclosing the challenge, from which positions estimates of the two dock bays and two racquetball target placards are determined based on the known geometry of the challenge. Next, the AMS approaches each of these 4 points of interest, orienting itself so the symbol is near the center of the camera's field of view. These images are inputted into the deep learning software to generate the best prediction of the symbol's shape and color. If this is the correct symbol, the docking or racquetball launching procedure begins. The docking procedure simply sets a new waypoint in the center of the bay, relying on the controller and obstacle avoidance system to reach this goal safely. For launching racquetballs, we found a more complicated feedback loop was required to account for wind, waves, and the drift of both the AMS

and the target. A quickly written and efficient computer vision script uses binary thresholding and edge detection to identify the black border around the targets at roughly 10 frames per second. This new position of the target is fed directly into the controller to make small movements to keep the AMS aligned. We intentionally bypass the obstacle avoidance system for these small adjustments for performance. The system constantly compares its real pose to the desired pose, only launching the racquetball when there is a low error.




Figure 8. Simple traditional computer vision to segment targets

IV. EXPERIMENTAL RESULTS

A. Simulator

The first phase for testing new software for NaviGator AMS is simulation. We use a modified version of VMRC, the beta platform for the virtual marine robotics challenge. This simulator uses similar technologies to modern 3D video games to render images for the simulated vision cameras and LIDAR. Every challenge present in RobotX 2018 is modeled in the simulator, allowing each task to be tested independently and in sequential runs similar to the finals of RobotX. Architecturally, the simulator uses Gazebo, an open source robotics simulator designed to integrate well with the ROS middleware we use. This allows us to run the exact same software in simulator as on the life platform, as the TCP socket interfaces for hardware (sensors and actuators) are fulfilled by the simulator. We added additional plugins to simulate the protocols of our student designed boards used for the emergency stop, pneumatic actuator, and passive sonar systems. The simulated hardware enables testing the integration of these systems into the higher level software without having physical access to the system. Simulation also makes the development of high level decision making programs, known as "missions", to proceed in parallel to perception software. Developers can optionally run the simulator in ground truth mode to receive perfect information about computer vision targets, nearby obstacles, and the position of the pinger. In this mode we can verify that the logic of the missions

is correct (the system moves correctly to complete the challenges) in ideal conditions. This separation of concerns allows the team to test a layer of our autonomy in isolation, which is essential for finding bugs and other design failures.



Figure 9. Simulated NaviGator ASV in a realistic environment

B. Field Testing

In addition to testing in the simulator, NaviGator ASV underwent significant lake testing. Over 130 hours of in-water testing were carried out in the form of day-long tests in the months leading up to the competition at a lake near UF. Lake testing offered real-life environmental factors that simulation cannot accurately provide, such as wind and current disturbances, various lighting conditions, and inclement weather.



Figure 10. Photo of NaviGator during lake testing

Field testing also offered a chance to test the mechanical systems of the ASV, such as actuators like the racquetball launcher, the strength of team-manufactured components, and the efficiency of the computer cooling system. The frequency and duration of testing helped to expose hardware failures that may have gone unnoticed until the competition. For example, the original sensor mast placed the Ubiquiti omnidirectional Wi-Fi antenna less than two inches away

from the Velodyne LIDAR. During field testing, the team found that the LIDAR was returning noisy data. However, when testing in the lab, the LIDAR data looked fine. Eventually the team determined that the only difference was that a wired connection was used to connect to the ASV while working in the lab, as opposed to the Wi-Fi connection that was used while field testing. It turns out that the Wi-Fi signal from the antenna was adding noise to the LIDAR data. Moving the Wi-Fi antenna further from the LIDAR solved the problem. This kind of issue would never have arisen during simulation. The detection of this and other flaws during testing prevented what would have been catastrophic failures during the competition. C. Field Element Construction In order to take full advantage of the realistic testing environment that the lake provides, field elements similar to those that will be used in the competition were constructed. The field elements were designed to be simple in construction and easy to deploy. Many of the elements were made of a PVC pipe frame that allowed for modular construction and easy assembly and disassembly. Buoyancy was provided by foam sheets and pool noodles fitted around the PVC pipes. The simplicity and light weight of the course elements allowed for quick and easy setup and teardown of the course using only a few team members in a kayak.

V. CONCLUSION

This paper presents the University of Florida's autonomous surface vehicle, NaviGator ASV, for use in the 2016 Maritime RobotX Challenge. Sacrificing speed for maneuverability, the vessel's four thrusters give the ASV an additional degree of freedom when compared to traditional skid-steer vessels. The novel use of an automotive cargo box for housing electronics created an open layout design that allowed for easy access and rapid repairs. An iterative approach and deep neural network pipeline created a strong software foundation that was exhaustively tested with over 130 hours of in-water testing. Team NaviGator AMS is ready for the 2018 Maritime RobotX Challenge due to extensively tested software, simple mechanical design, and robust electronics.

VI. ACKNOWLEDGEMENT

Team NaviGator AMS would like to acknowledge everyone who has supported the team throughout the year, including the University of Florida's Herbert Wertheim College of Engineering, the Electrical and Computer Engineering department, the Mechanical and Aerospace Engineering department, as well as the labs of MIL and CIMAR. The team would like to extend an appreciative thank you to their advisers: Dr. Eric Schwartz, Dr. Carl Crane, Dr. Ira Hill, and Shannon Ridgeway. The latest Team NaviGator AMS developments can be found at www.NaviGatorUF.org.

REFERENCES

- [1] S. M. Lavalle, “Rapidly-Exploring Random Trees: A New Tool for Path Planning,” 1998.
- [2] A. Devarakonda, K. Fountoulakis, J. Demmel, and M. W. Mahoney, “Avoiding communication in primal and dual block coordinate descent methods,” *arXiv:1612.04003 [cs]*, Dec. 2016.
- [3] “Adaptive control,” *Wikipedia*, Nov. 2018.